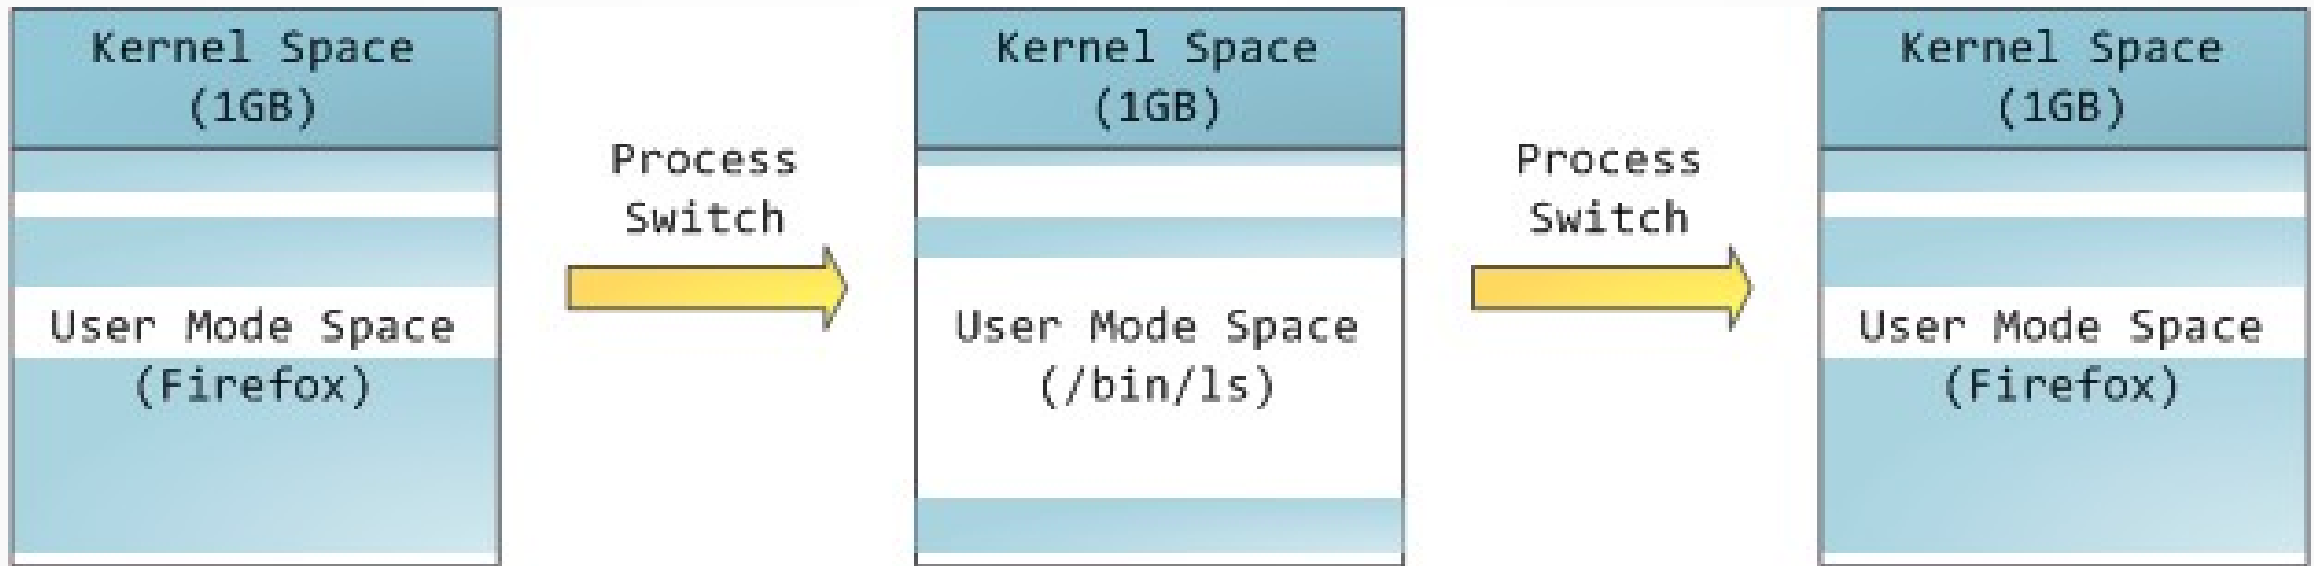


Memory management in Linux

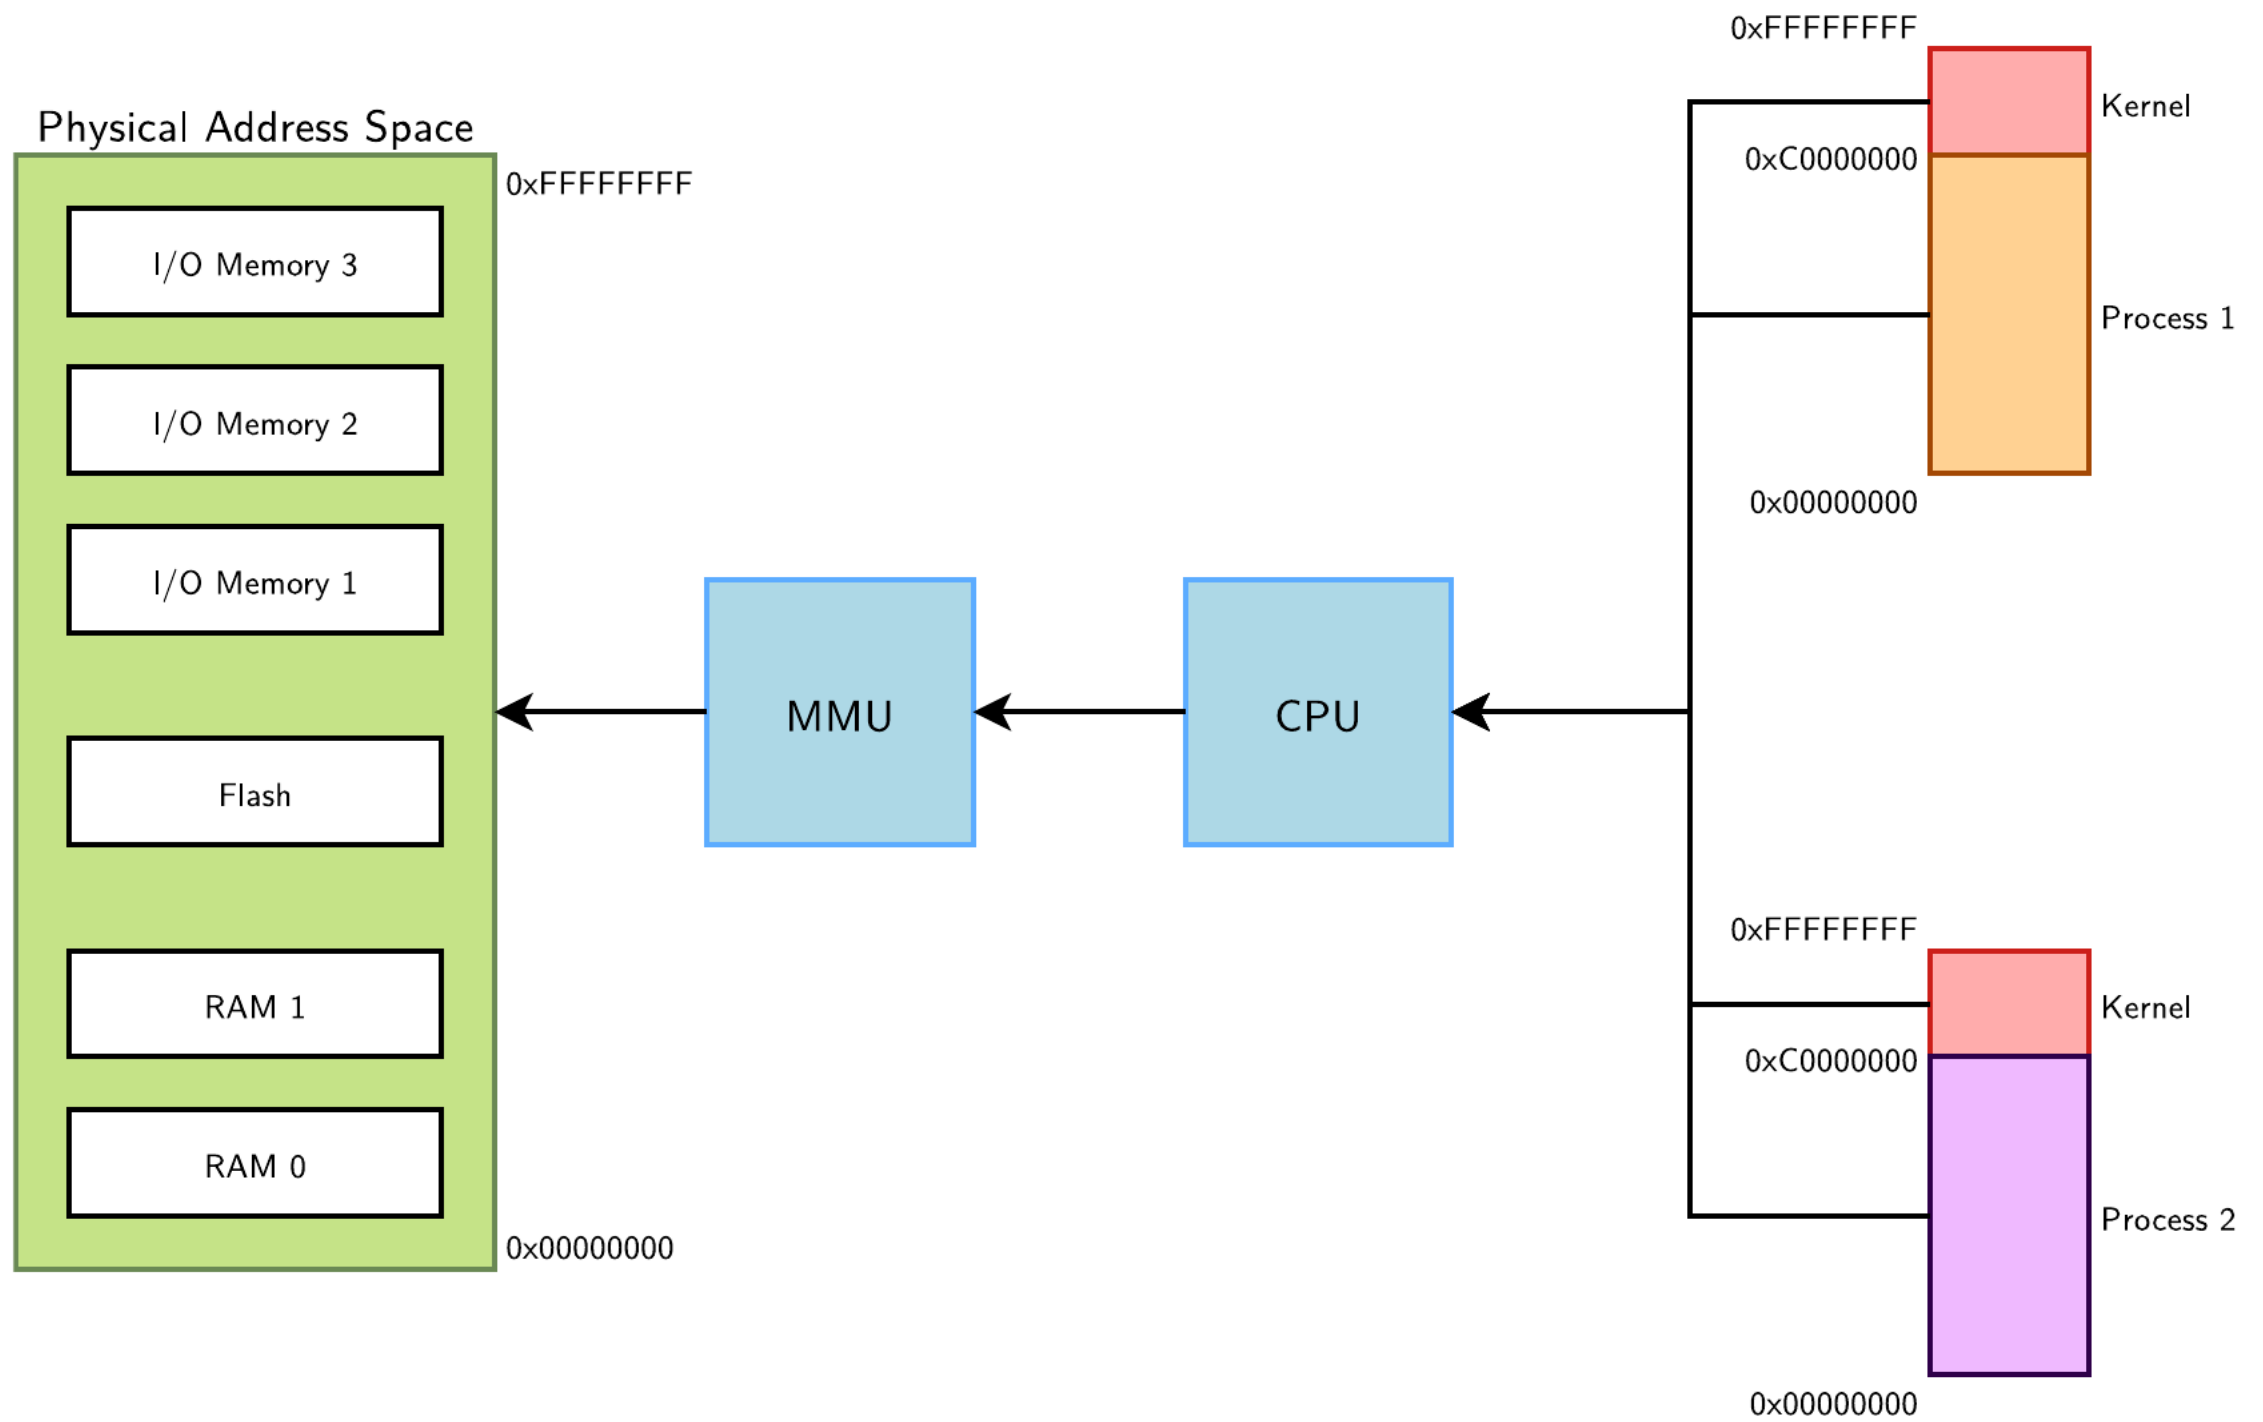
Memory management tasks

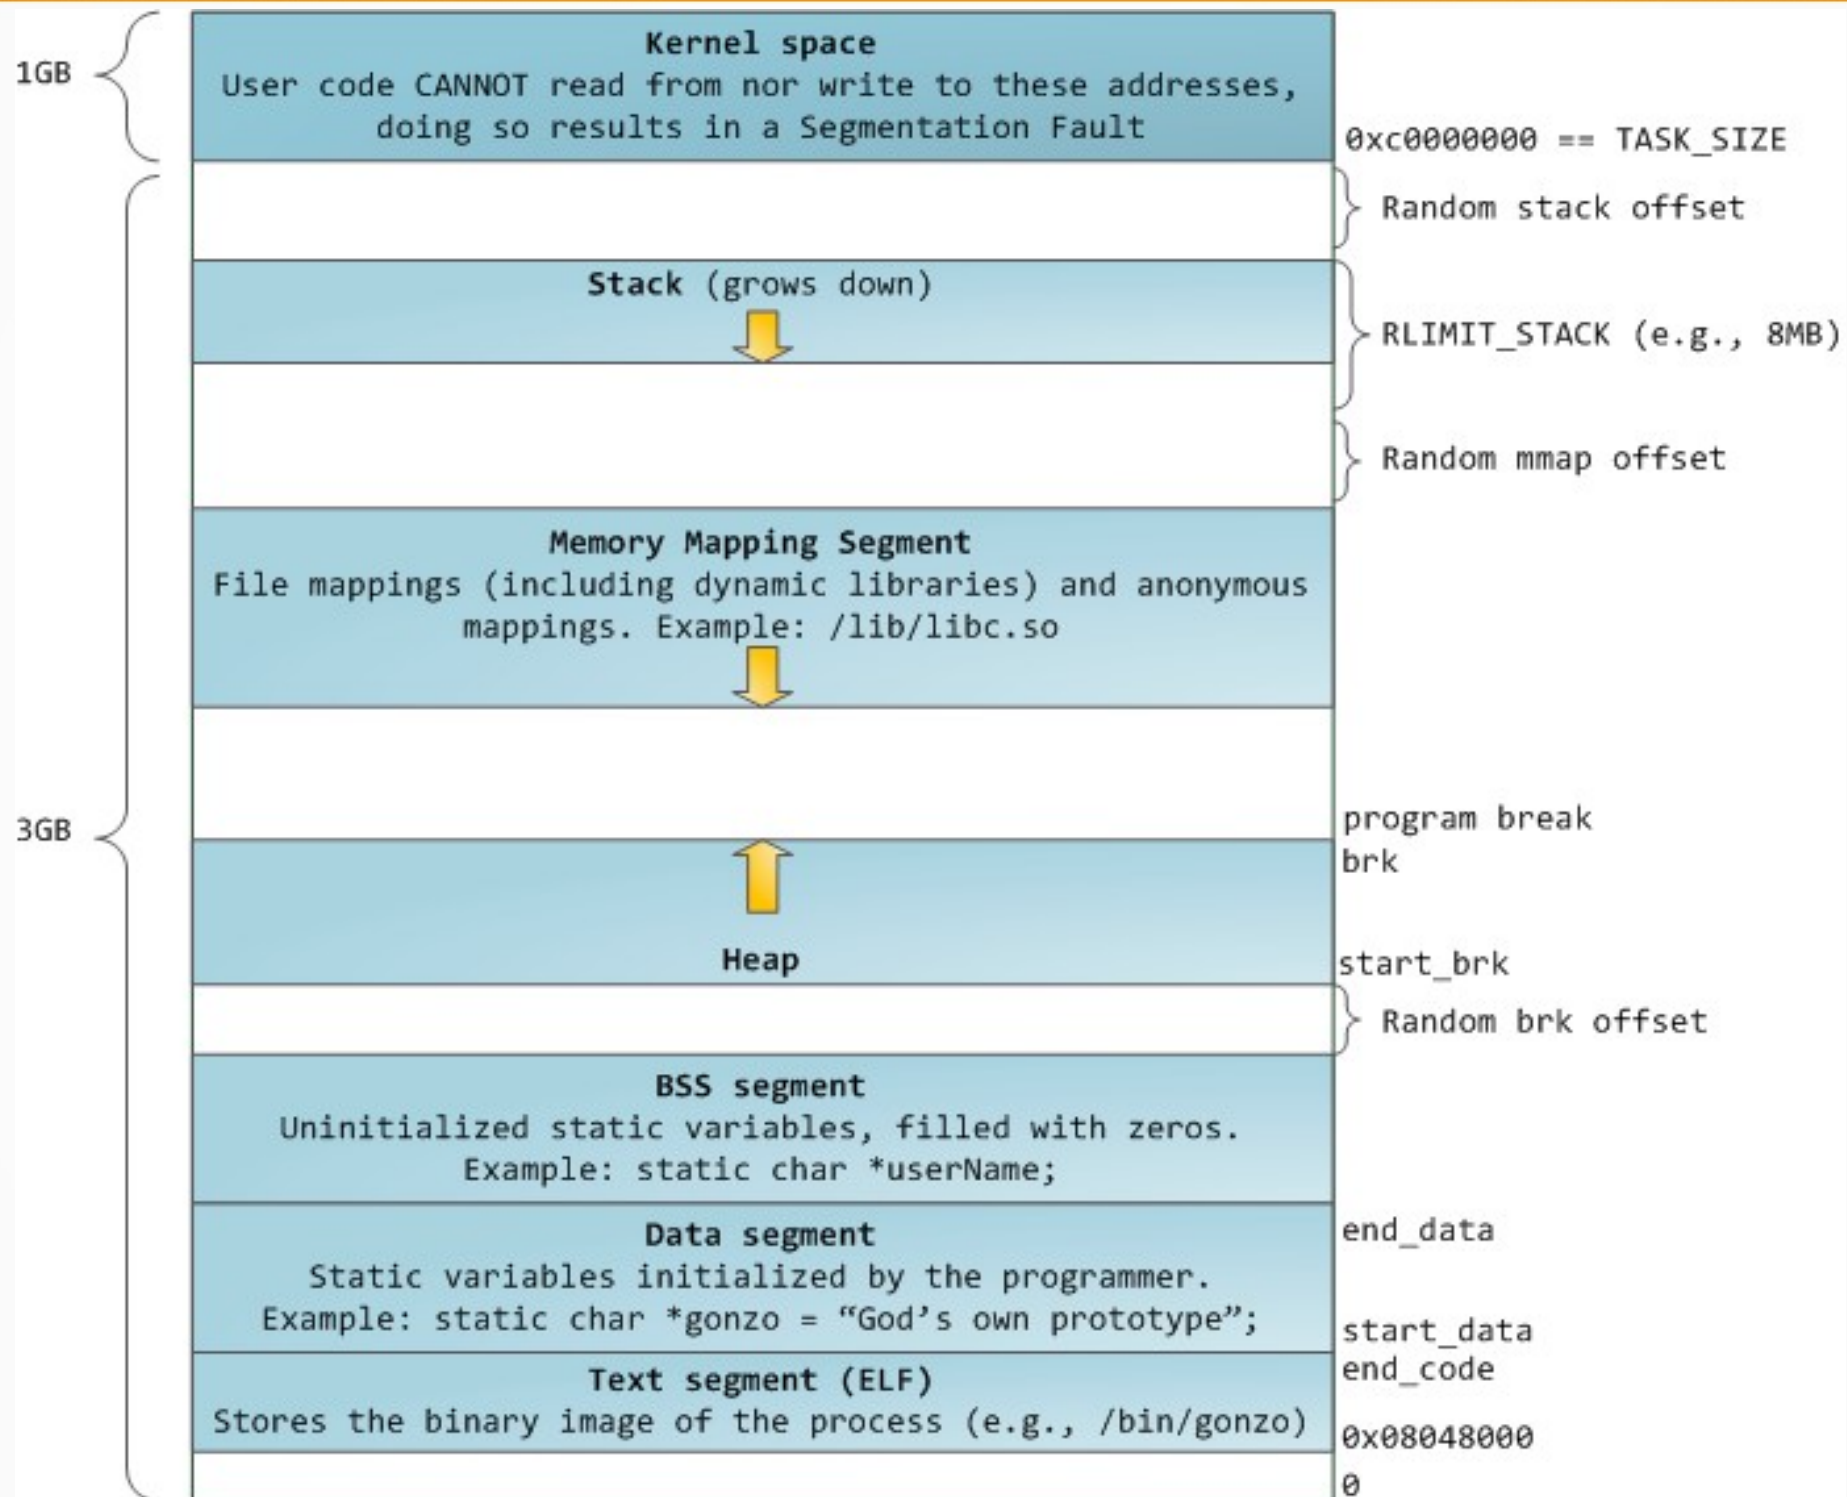
- Physical memory management
- Virtual memory allocation
- Physical / virtual memory mapping, PTE management
- Memory allocation for kernel needs

Virtual memory organization



- Kernel code and data are always addressable, ready to handle interrupts or system calls at any time
- Mapping for the user-mode portion of the address space changes whenever a process switch happens





Stack

- If a process pushes data beyond stack size, the CPU will trigger a page fault.
- The page fault handler detects the address is just beyond the stack, and allocates a new page to extend the stack.
- See `__do_page_fault()` in `arch/{x86,arm,arm64...}/mm/fault.c`, then `expand_stack()` and `acct_stack_growth()`
- If the stack size is more than `RLIMIT_STACK`, a `SIGSEGV` signal is generated (segmentation fault).

mmap()

- `mmap()` is the standard way to allocate large amounts of memory from user space.
- Can map contents of files directly to memory, which is used for example for loading dynamic libraries.
- `MAP_ANONYMOUS` flag causes `mmap()` to allocate normal memory for the process. The `MAP_SHARED` flag can make the allocated pages shareable with other processes.
- If you request more than `M_MMAP_THRESHOLD` (128 kB by default, adjustable via `mallopt()`) via `malloc()`, anonymous mapping is used instead of heap memory.

Heap

- Allocating from heap is the standard way to allocate small amounts of memory from user space: `malloc()` and friends in C, 'new' keyword in C++, etc.
- If there is enough space in the heap to satisfy a memory request, it can be handled by the language runtime without kernel involvement. Otherwise the heap is enlarged via `sbrk()/brk()` (see `do_brk()` in `mm/mmap.c`).
- Large allocations use `mmap()` instead of `brk()`.
- May become fragmented eventually:



/proc/<pid>/maps

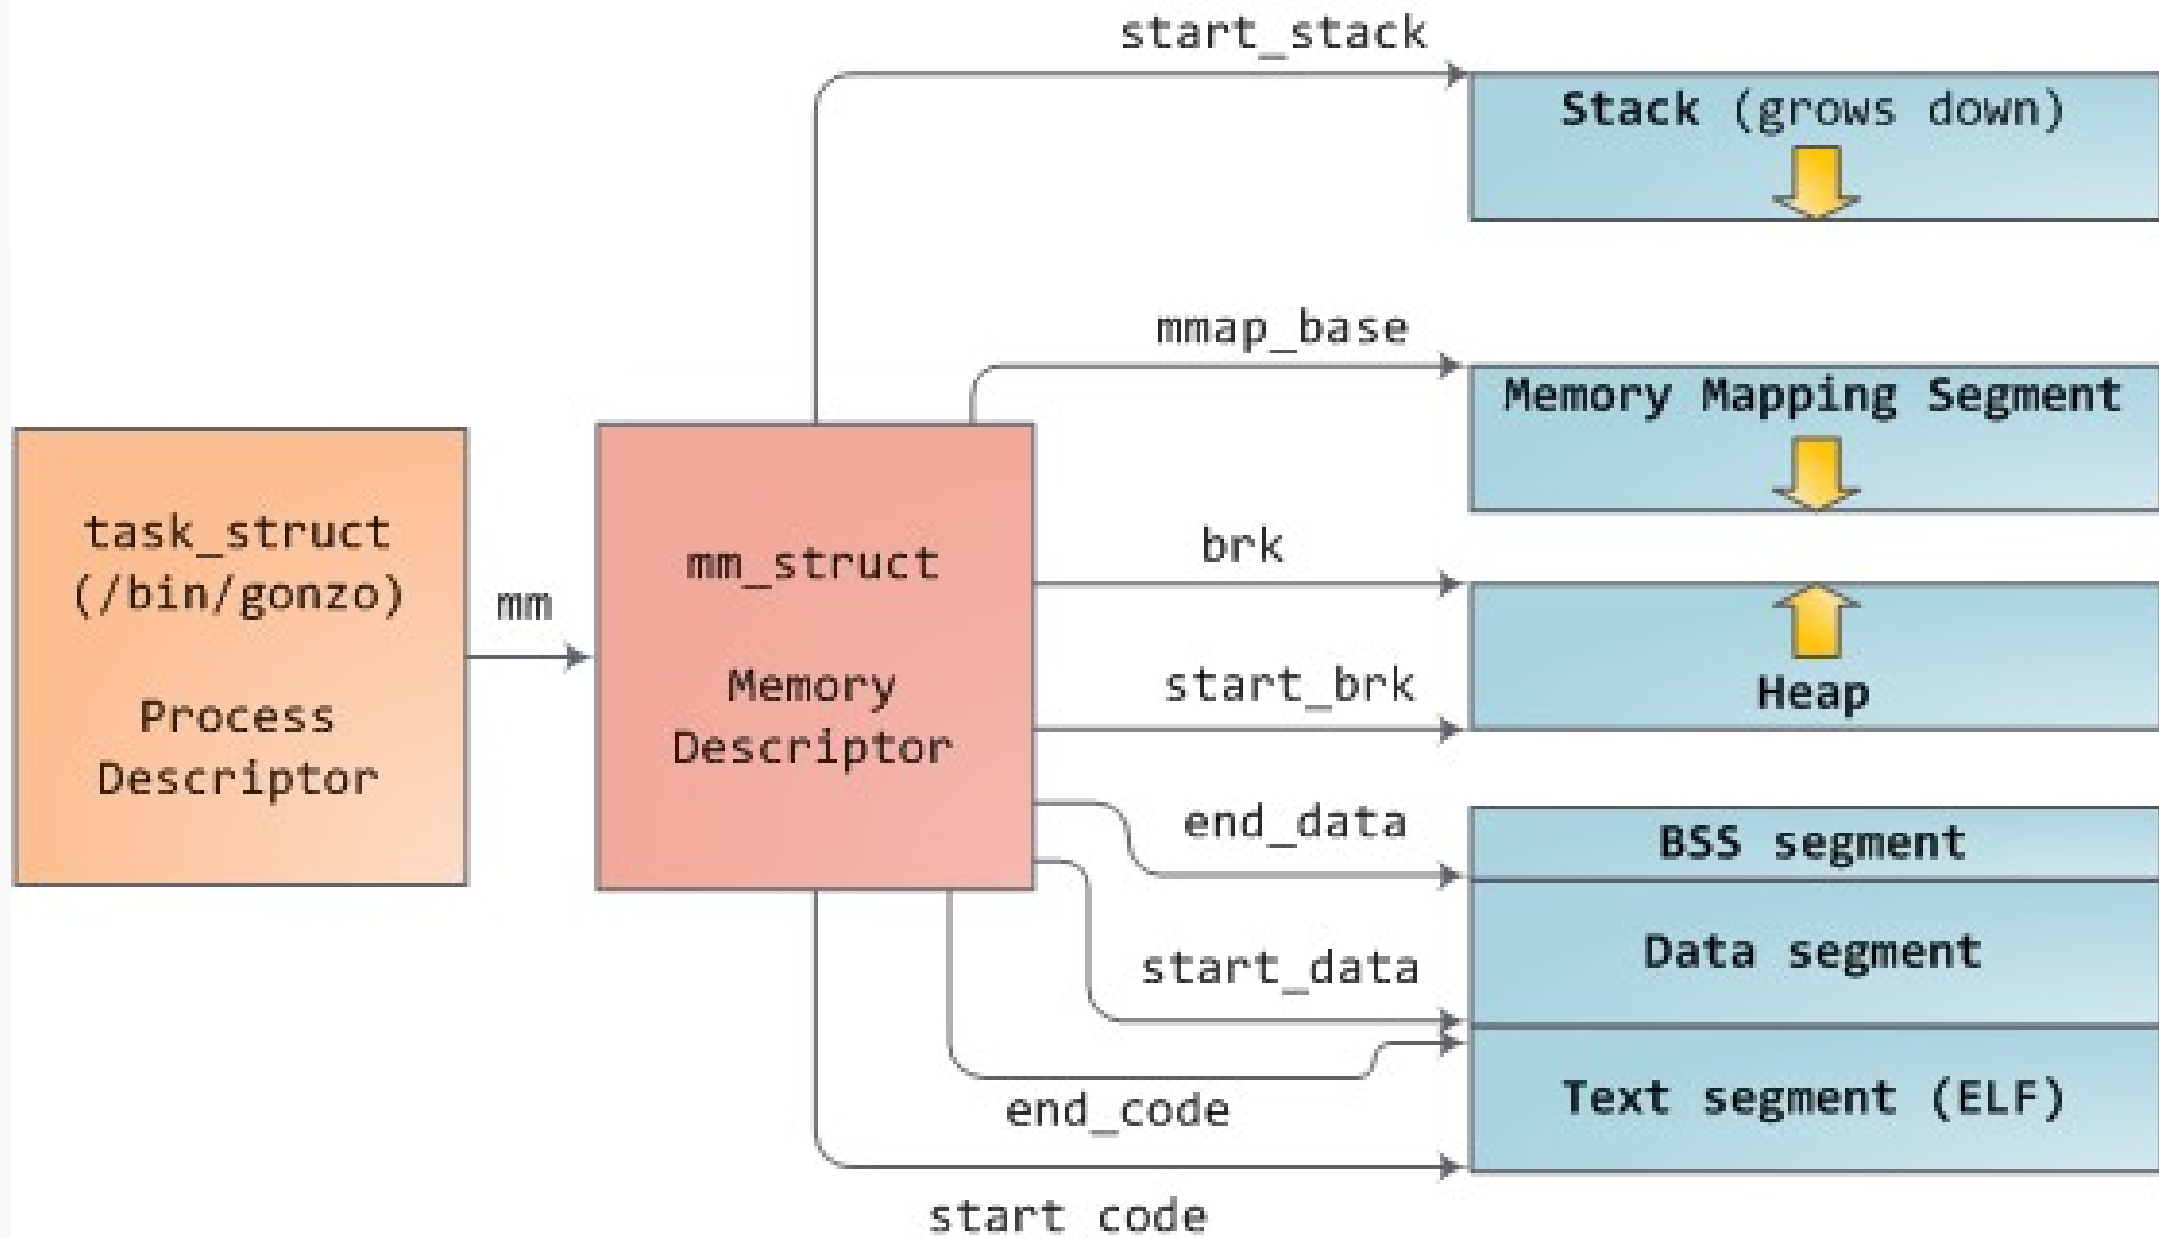
- The /proc/PID/maps file containing the currently mapped memory regions and their access permissions.

The format is:

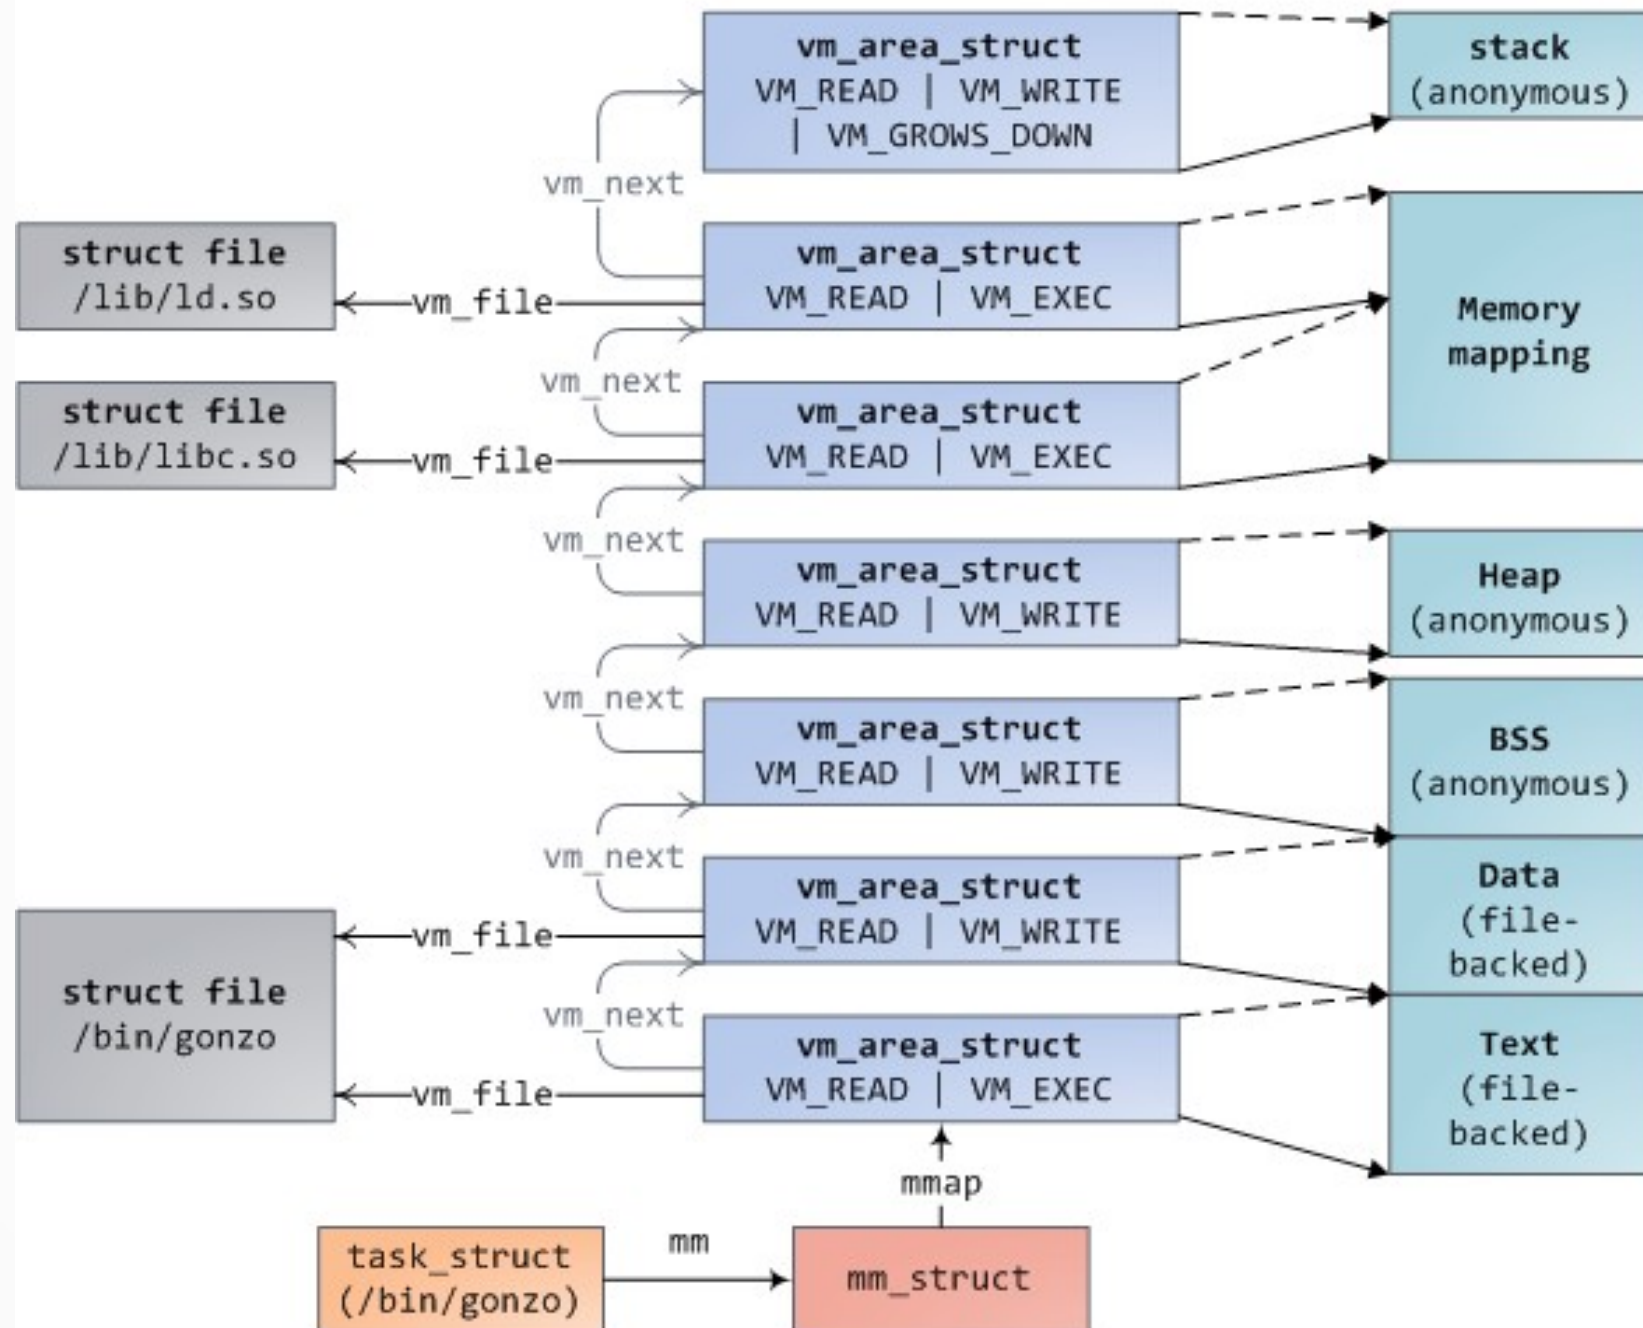
address	perms	offset	dev	inode	pathname
~/ cat /proc/self/maps					
55de9ca91000-55de9ca99000	r-xp	00000000	fc:00	23855162	/bin/cat
55de9cc98000-55de9cc99000	r--p	00007000	fc:00	23855162	/bin/cat
55de9cc99000-55de9cc9a000	rw-p	00008000	fc:00	23855162	/bin/cat
55de9df89000-55de9dfaa000	rw-p	00000000	00:00	0	[heap]
7f32494c1000-7f32497da000	r--p	00000000	fc:00	5902488	/usr/lib64/locale/locale-archive
7f32497da000-7f3249993000	r-xp	00000000	fc:00	25668388	/lib64/libc-2.26.so
7f3249993000-7f3249b92000	---p	001b9000	fc:00	25668388	/lib64/libc-2.26.so
7f3249b92000-7f3249b96000	r--p	001b8000	fc:00	25668388	/lib64/libc-2.26.so
7f3249b96000-7f3249b98000	rw-p	001bc000	fc:00	25668388	/lib64/libc-2.26.so
7f3249b98000-7f3249b9c000	rw-p	00000000	00:00	0	
7f3249b9c000-7f3249bc1000	r-xp	00000000	fc:00	25669421	/lib64/ld-2.26.so
7f3249d81000-7f3249d83000	rw-p	00000000	00:00	0	
7f3249d9e000-7f3249dc0000	rw-p	00000000	00:00	0	
7f3249dc0000-7f3249dc1000	r--p	00024000	fc:00	25669421	/lib64/ld-2.26.so
7f3249dc1000-7f3249dc2000	rw-p	00025000	fc:00	25669421	/lib64/ld-2.26.so
7f3249dc2000-7f3249dc3000	rw-p	00000000	00:00	0	
7ffe92c6e000-7ffe92c8f000	rw-p	00000000	00:00	0	[stack]
7ffe92cf3000-7ffe92cf6000	r--p	00000000	00:00	0	[vvar]
7ffe92cf6000-7ffe92cf8000	r-xp	00000000	00:00	0	[vdso]
ffffffffffff600000-ffffffffffff601000	r-xp	00000000	00:00	0	[vsyscall]

See <https://www.kernel.org/doc/Documentation/filesystems/proc.txt>

Managing virtual memory. mm_struct

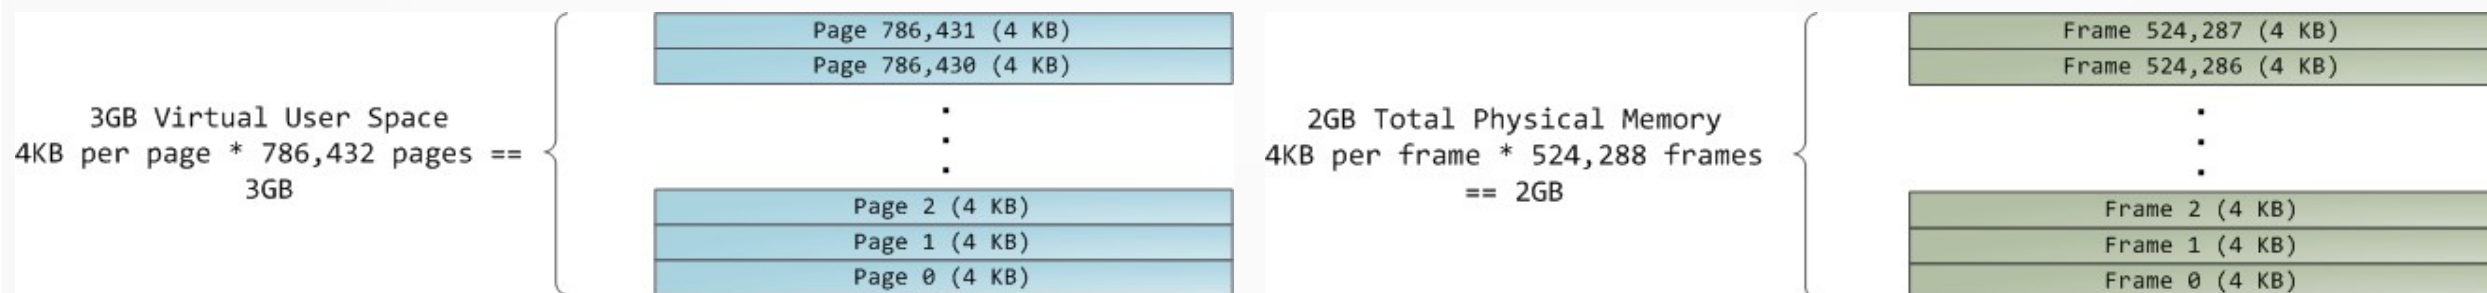


-----> vm_end: first address **outside** virtual memory area
-----> vm_start: first address **within** virtual memory area



Translation of virtual memory to physical memory

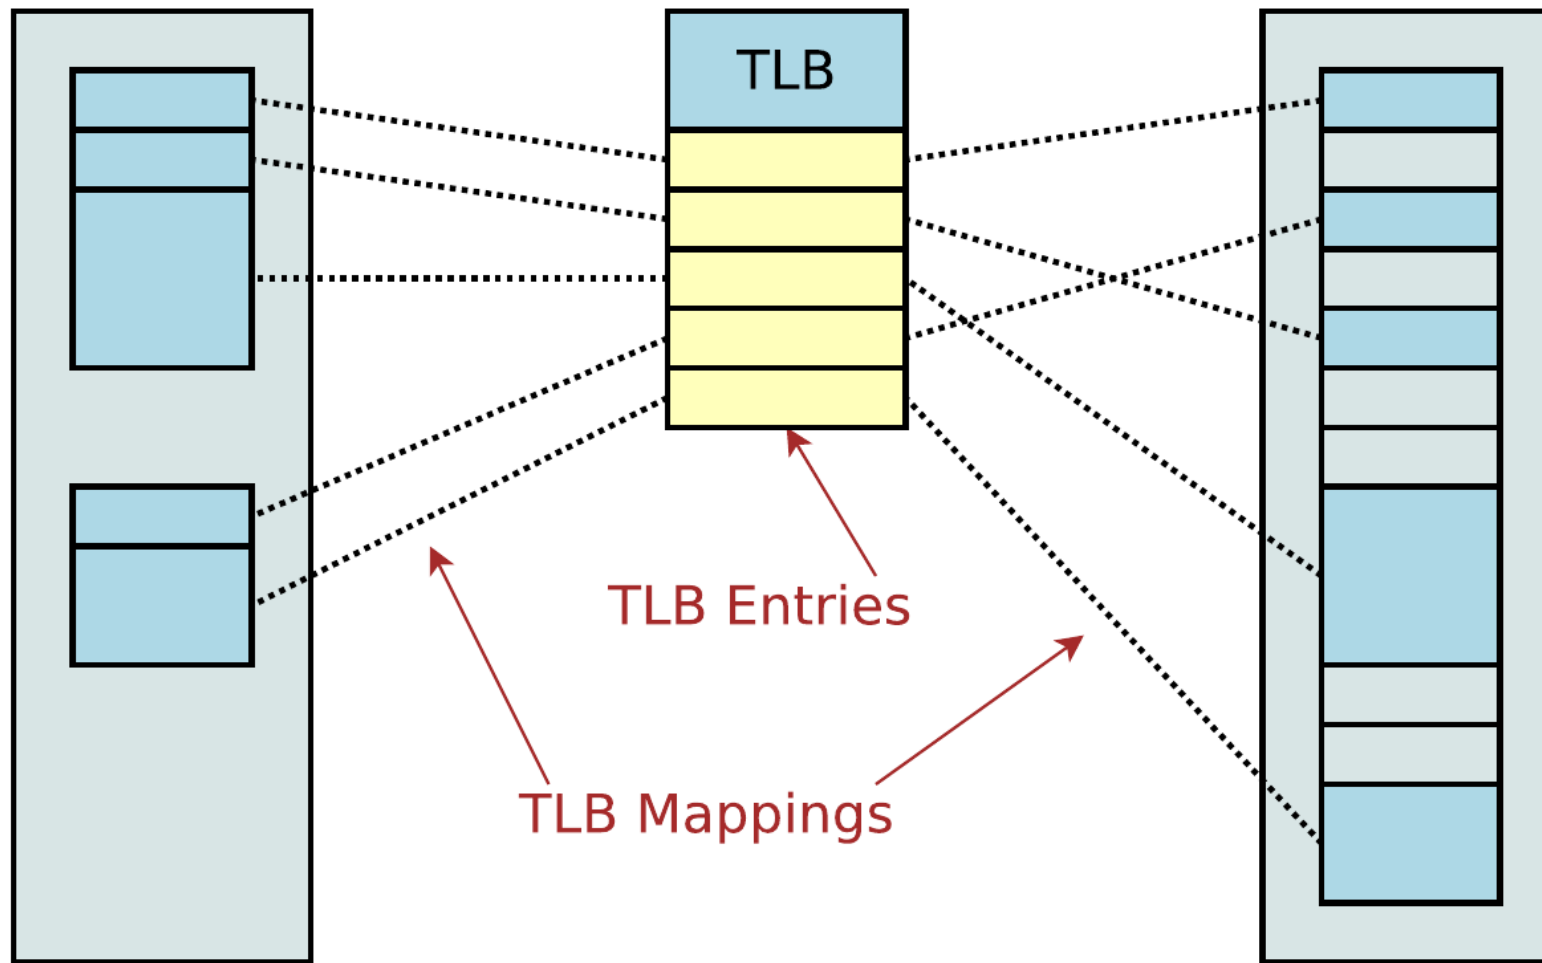
- The 4GB virtual address space is divided into pages (4 KB by default, 2 MB or 1 GB is also supported on x86_64).
- The size of a VMA must be a multiple of page size.
- MMU uses TLB or page tables to translate a virtual address into a physical memory address.



TLB mappings

User Virtual Address Space

Physical Address Space

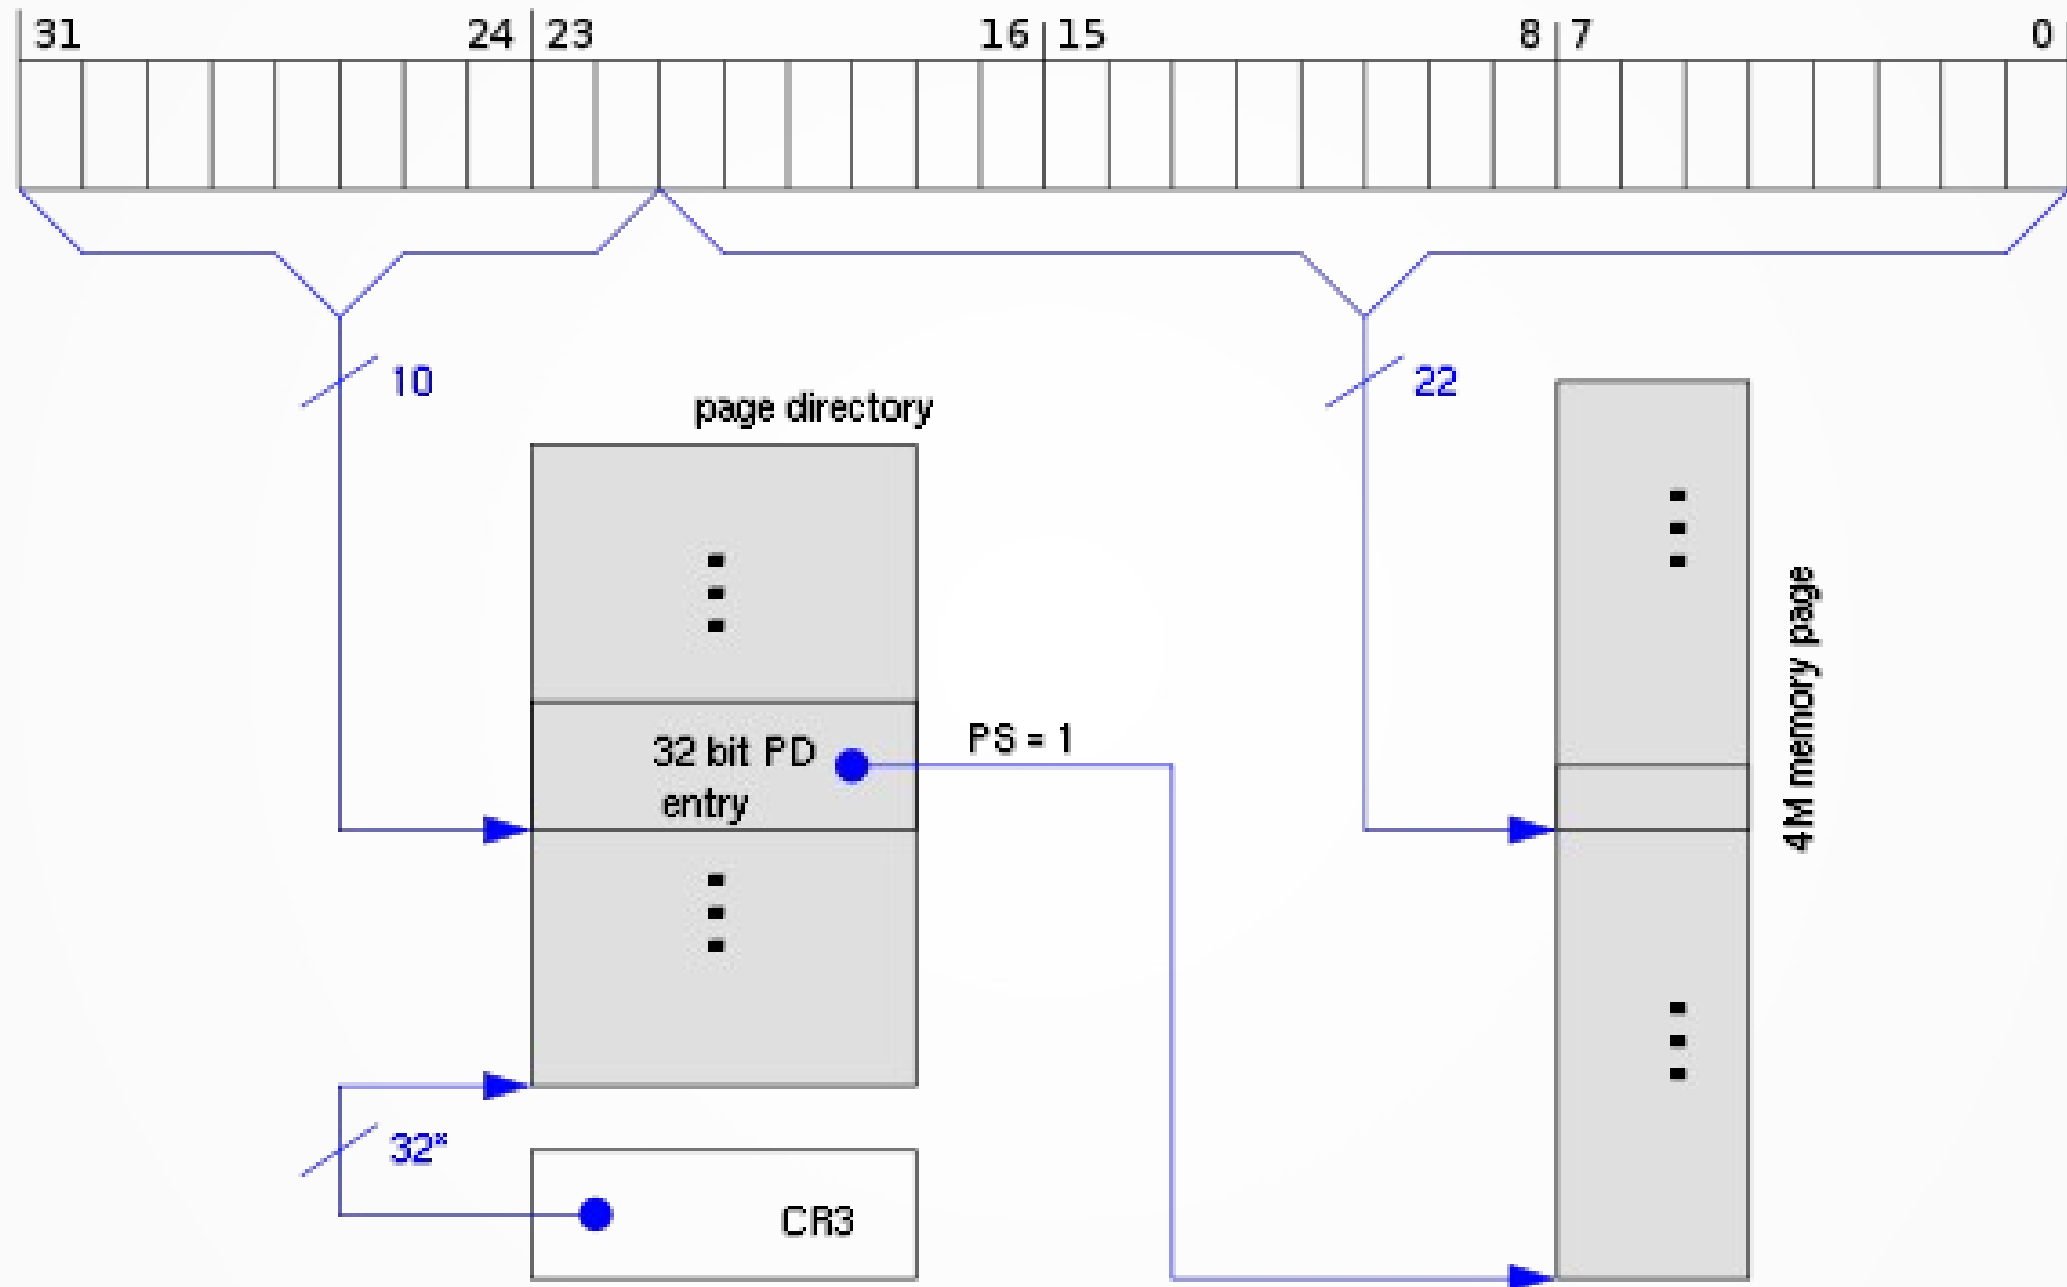


- Context switch loads new PDG pointer to CR3 and flushes TLB.

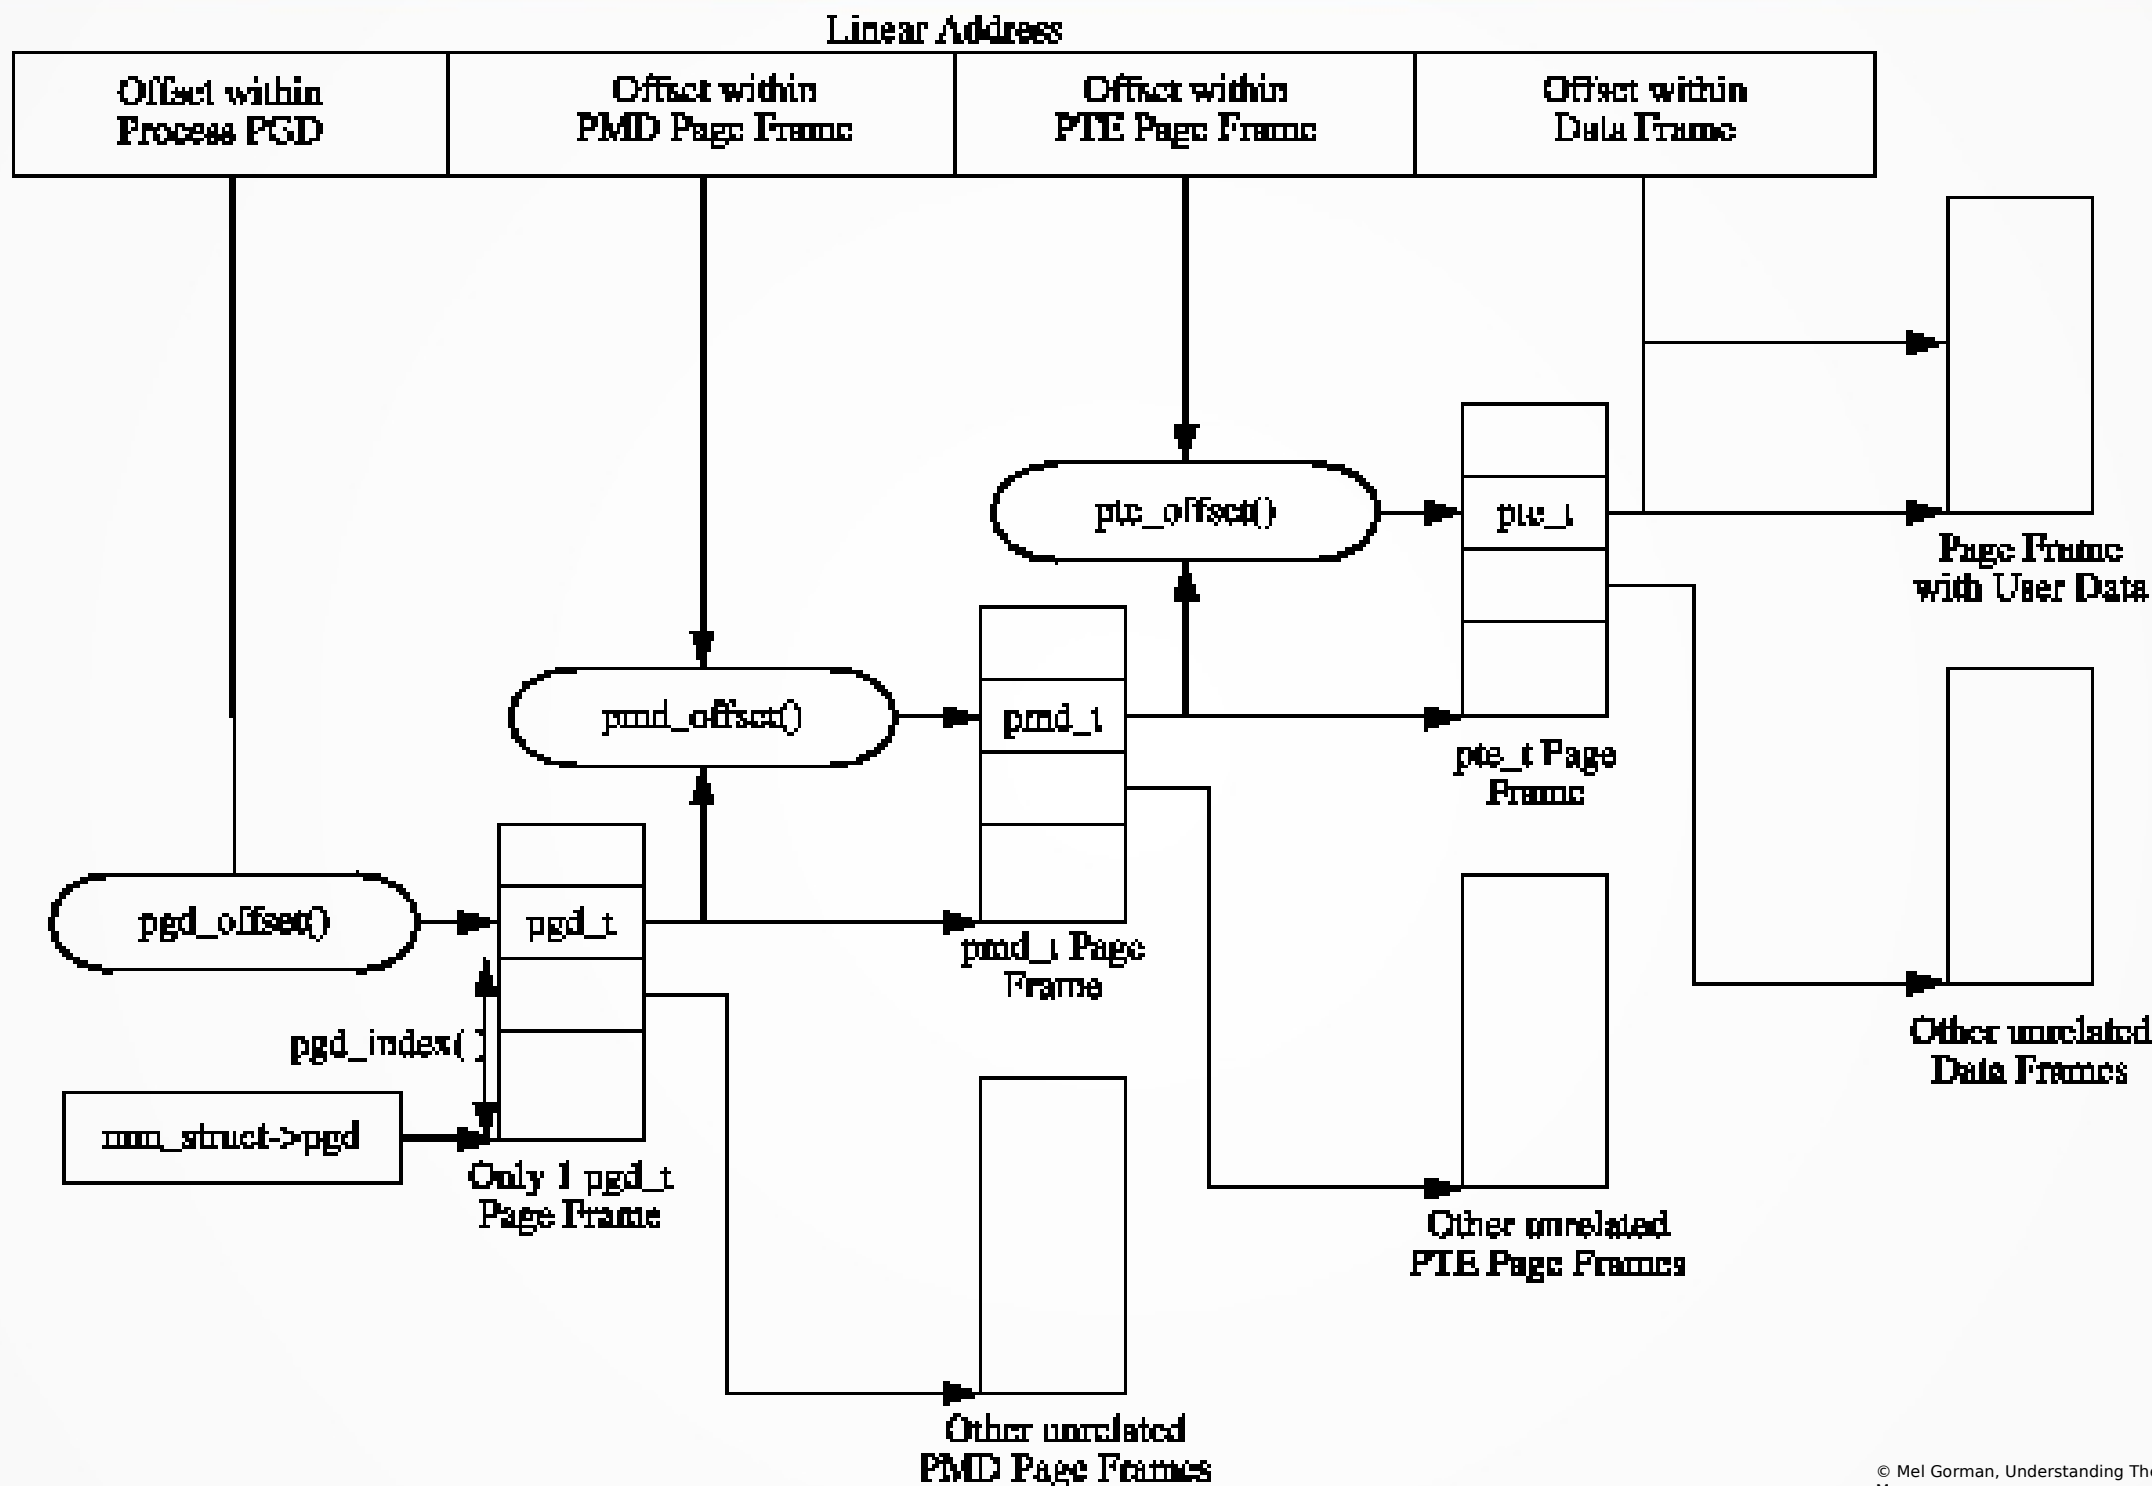
Page tables

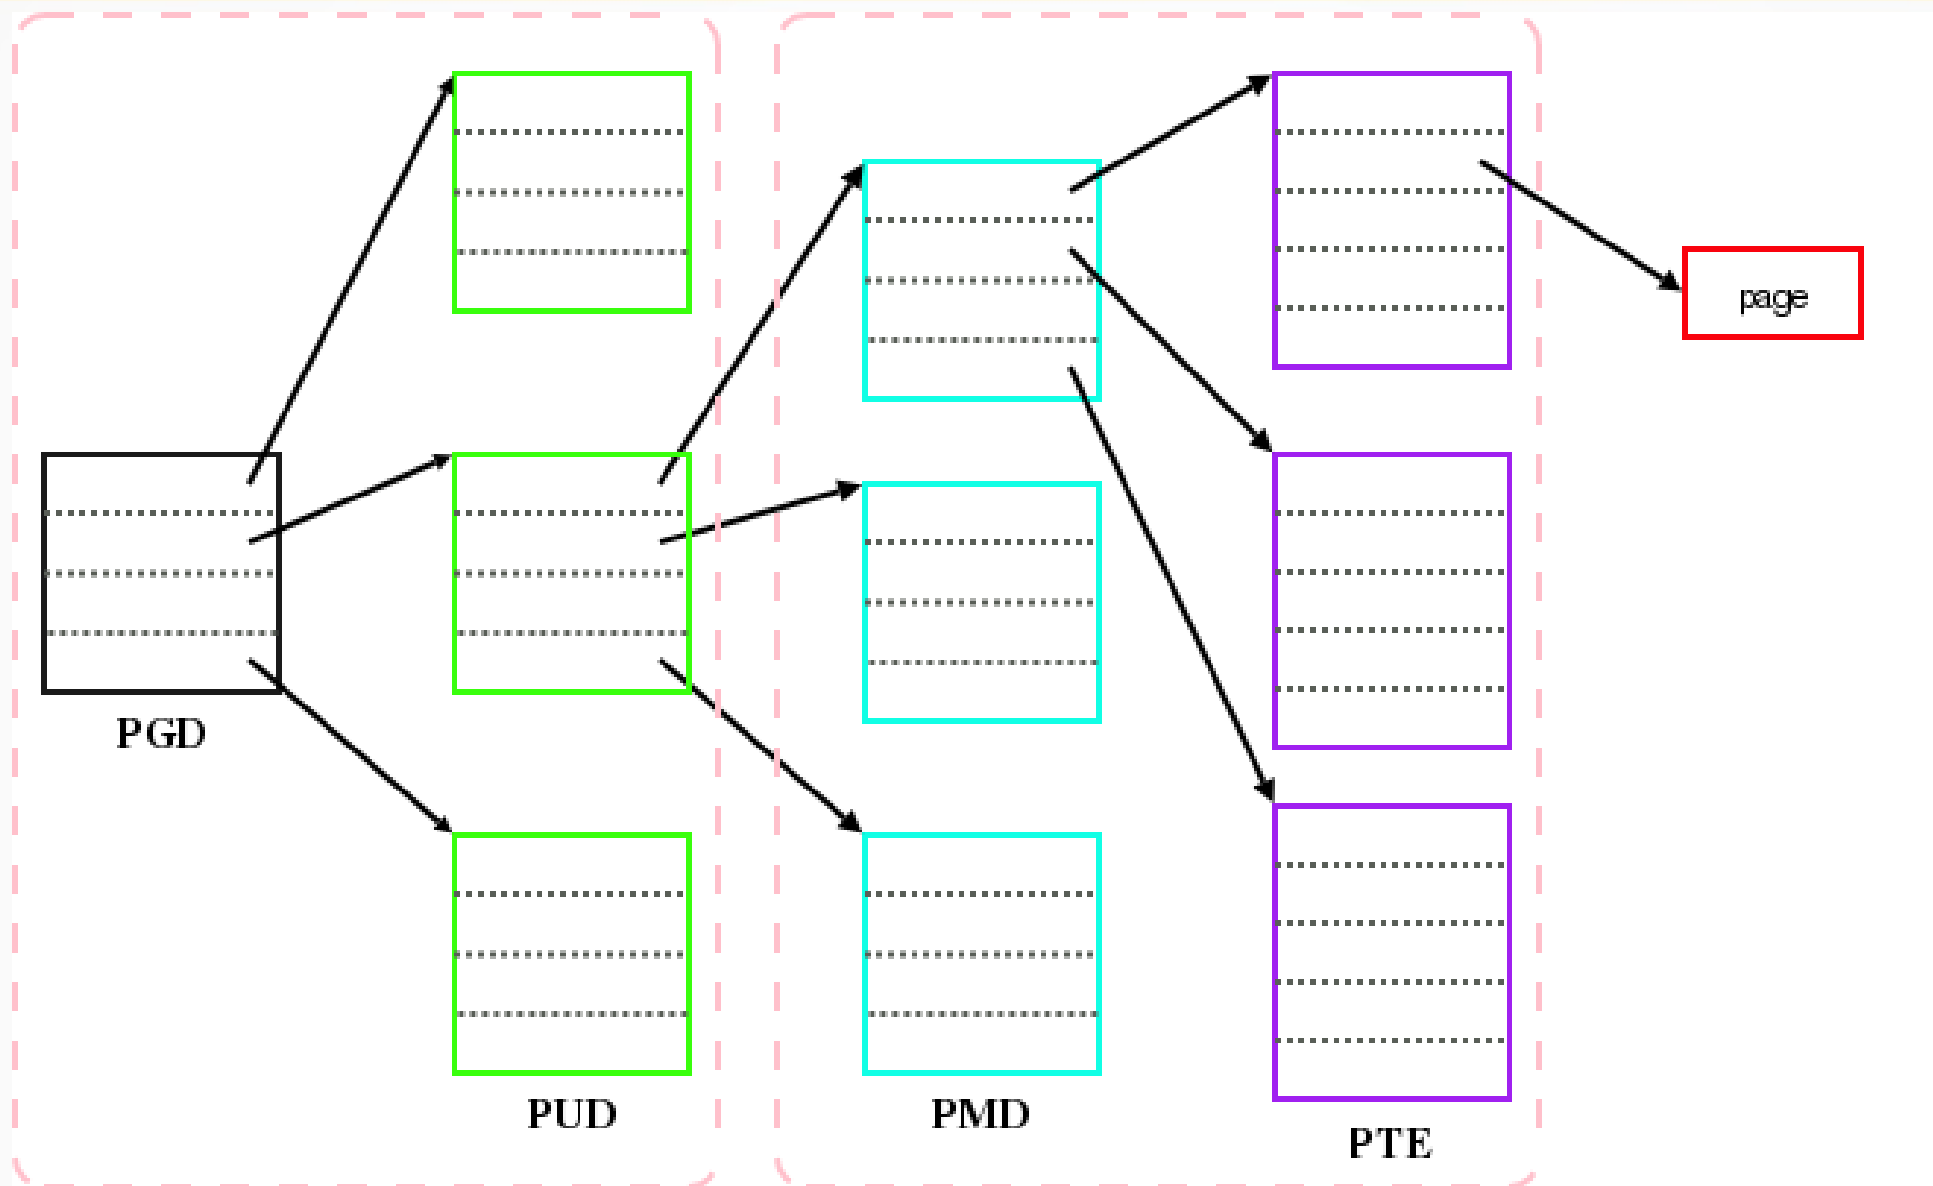
- Each process has a pointer (`mm_struct->pgd`) to its own Page Global Directory (PGD) which is a physical page frame.
- The page tables are loaded differently depending on the architecture. On the x86, the process page table is loaded by copying `mm_struct->pgd` into the `cr3` register.

Linear address:



*) 32 bits aligned to a 4-KByte boundy





Architecture	Bits used			
	PGD	PUD	PMD	PTE
i386	22-31			12-21
i386 (PAE mode)	30-31		21-29	12-20
x86-64	39-46	30-38	21-29	12-20

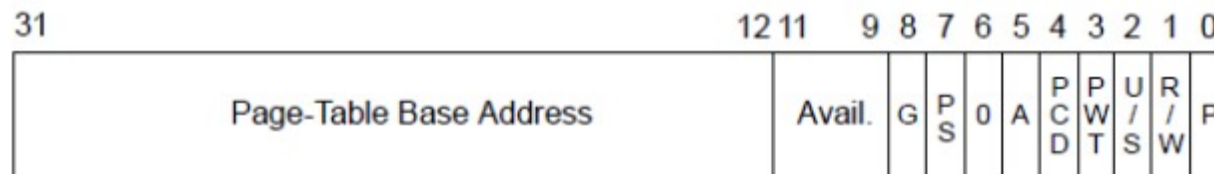
ARM: it's complicated...

See https://elinux.org/Tims_Notes_on_ARM_memory_allocation

Page tables

- Page directory and page table entries:

Page-Directory Entry (4-KByte Page Table)



Available for system programmer's use _____

Global page (Ignored) _____

Page size (0 indicates 4 KBytes) _____

Reserved (set to 0) _____

Accessed _____

Cache disabled _____

Write-through _____

User/Supervisor _____

Read/Write _____

Present _____

Page-Table Entry (4-KByte Page)



Available for system programmer's use _____

Global page _____

Reserved (set to 0) _____

Dirty _____

Accessed _____

Cache disabled _____

Write-through _____

User/Supervisor _____

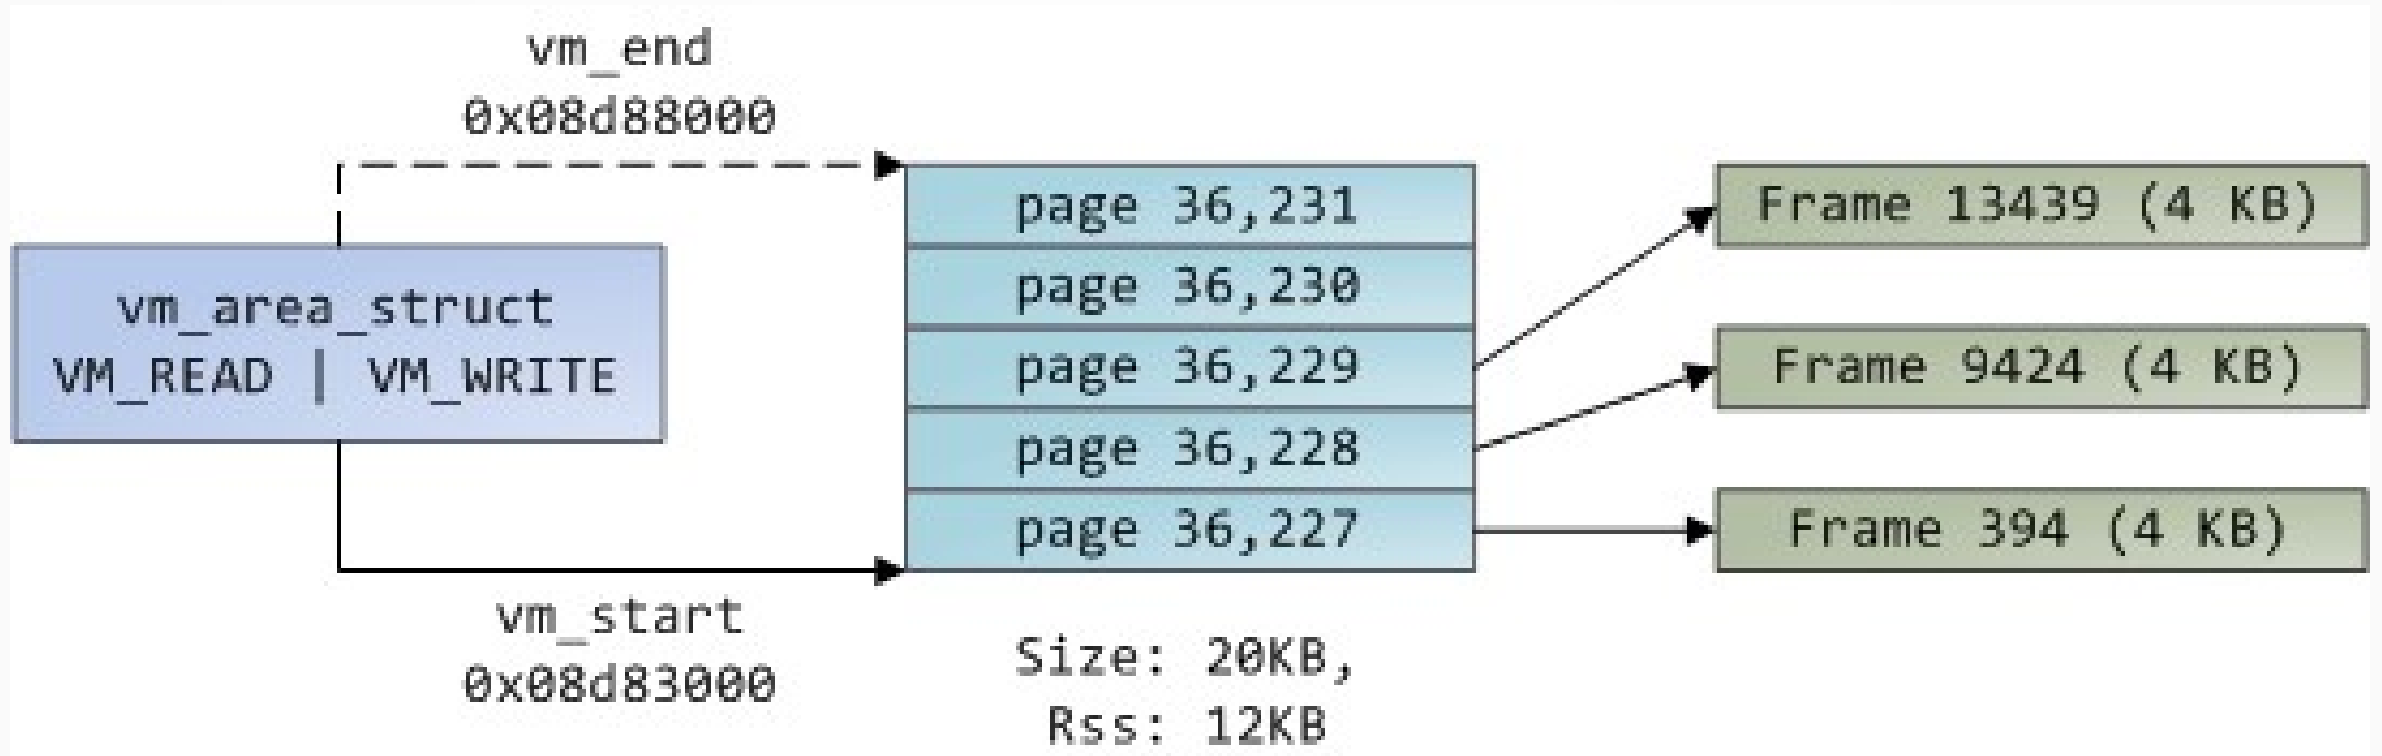
Read/Write _____

Present _____

Page frames

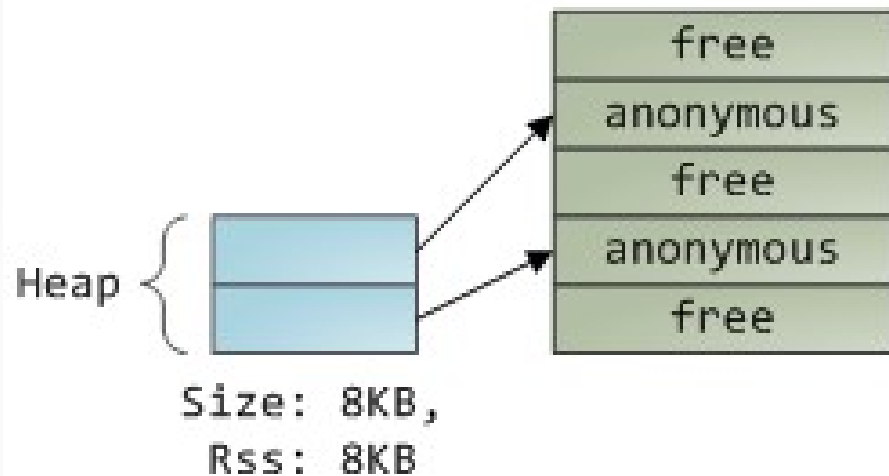
- In Linux each page frame is tracked by a descriptor and several flags.
- Physical memory is managed with the buddy memory allocator.
- A page frame is free if it's available for allocation via the buddy system.
- An allocated page frame might be **anonymous** (holding program data), or **page cache** (holding data stored in a file or block device).

Arrows represent page table entries mapping pages onto page frames

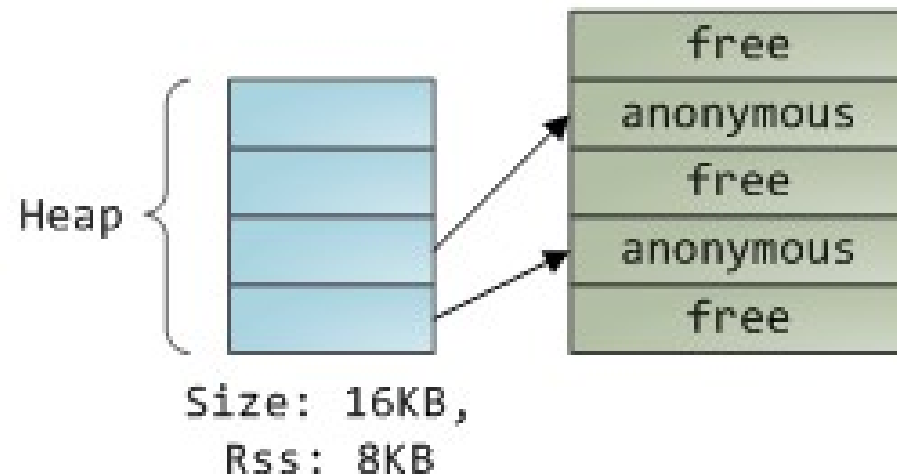


Some virtual pages lack arrows; this means their corresponding PTEs have the **Present** flag clear

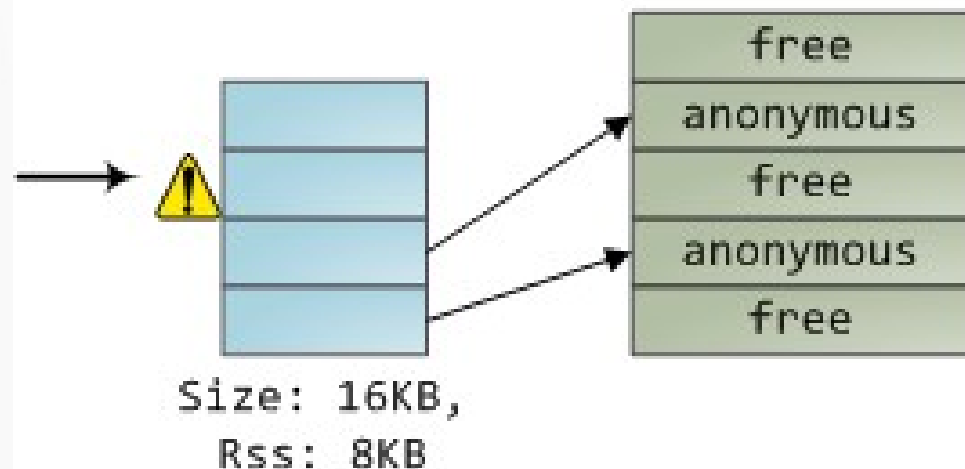
1. Program calls `brk()` to grow its heap



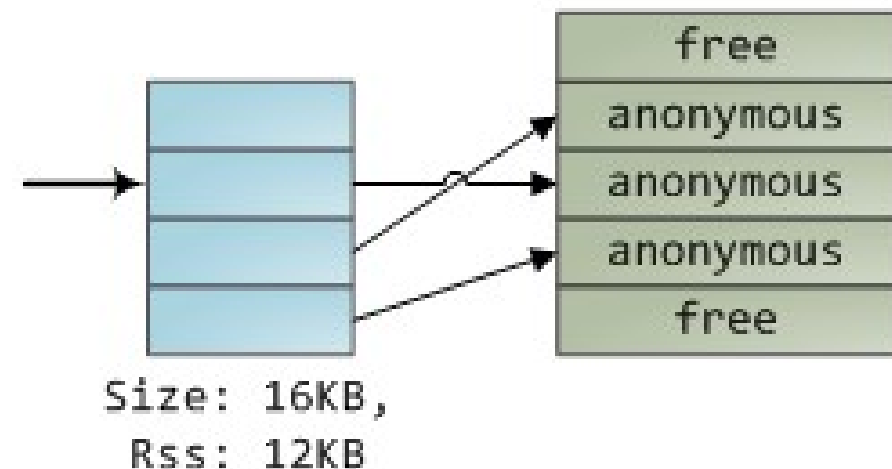
2. `brk()` enlarges heap VMA.
New pages are **not** mapped onto physical memory.



3. Program tries to access new memory.
Processor page faults.



4. Kernel assigns page frame to process,
creates PTE, resumes execution. Program is
unaware anything happened.

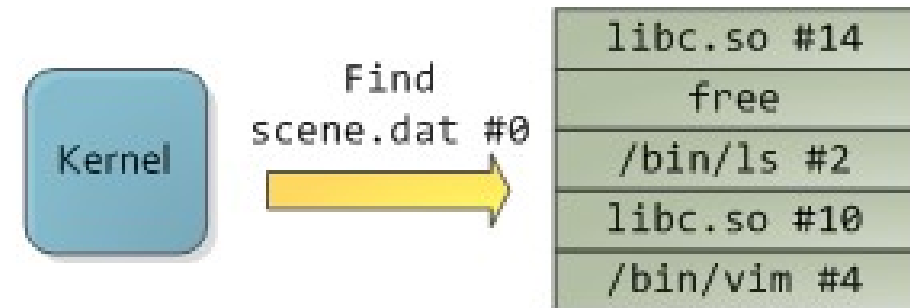


Page cache

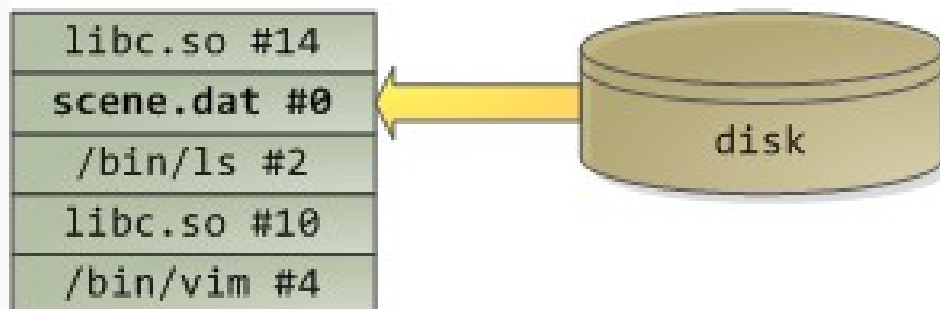
1. Render asks for 512 bytes of scene.dat starting at offset 0.



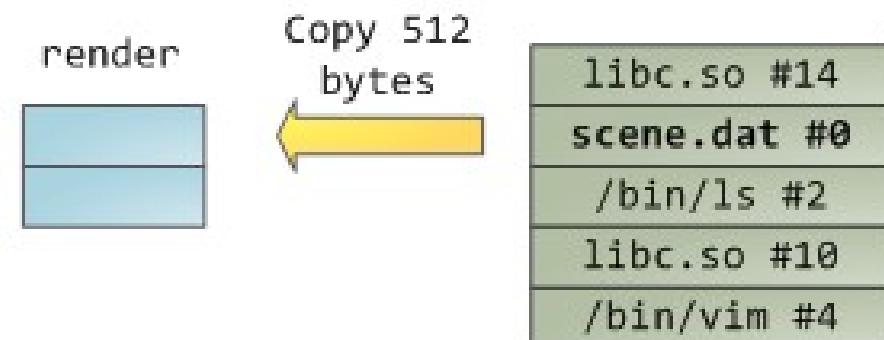
2. Kernel searches the page cache for the 4KB chunk of scene.dat satisfying the request. Suppose the data is not cached.



3. Kernel allocates page frame, initiates I/O requests for 4KB of scene.dat starting at offset 0 to be copied to allocated page frame

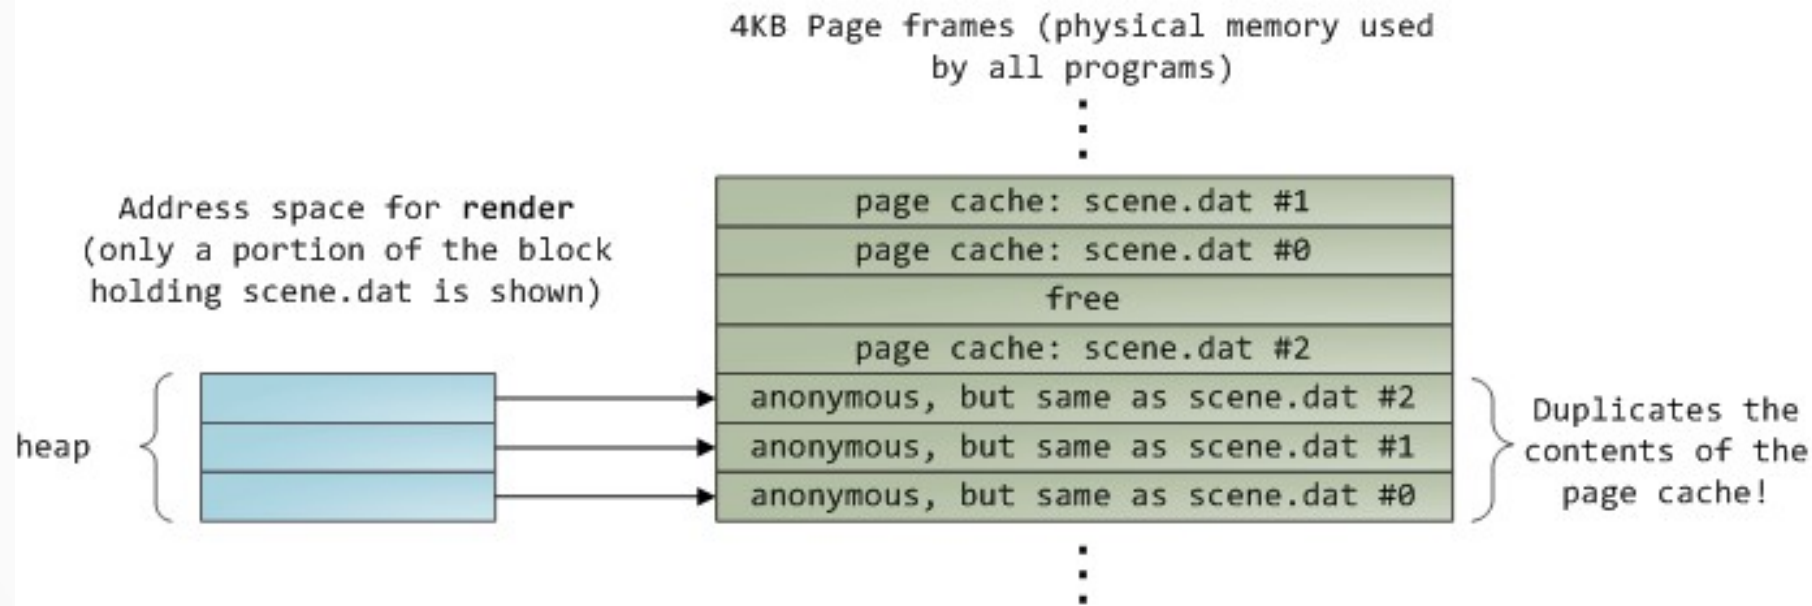


4. Kernel copies the requested 512 bytes from page cache to user buffer, read() system call ends.

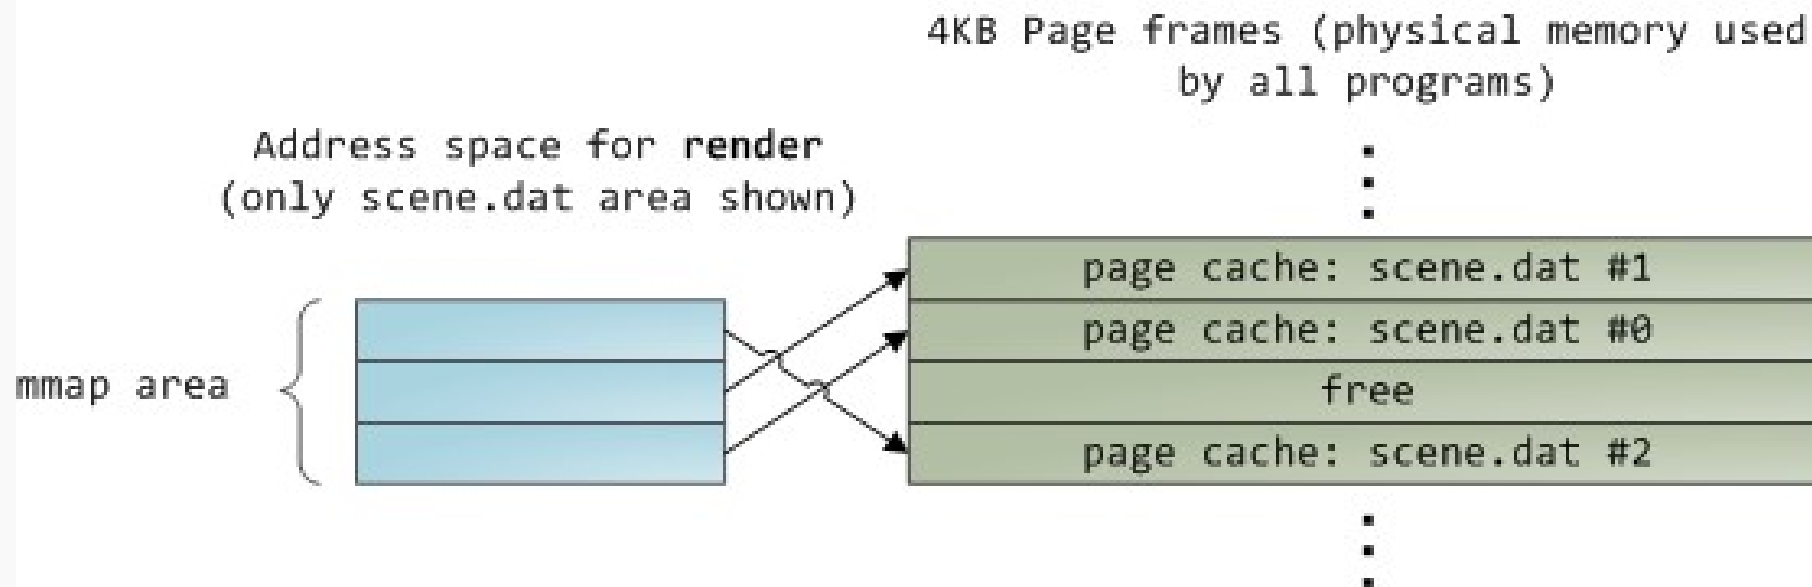


Page cache and memory-mapped files

read():



mmap():



Freeing/reclaiming page-cached memory

- When a process is terminated, the pages storing the mapped file are not freed.
- As long as there's enough free physical memory, the page cache won't be freed.
- It is not dependent on a particular process, page cache is a system-wide resource.

htop:

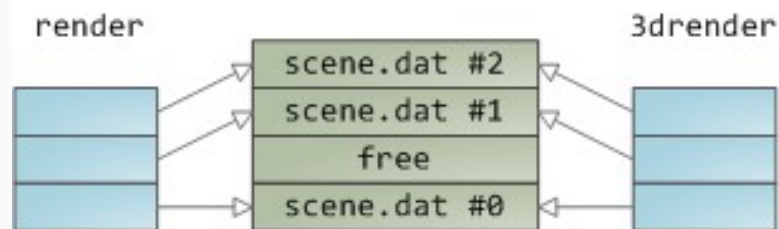
Mem[5.06G/11.6G]						
~/ free -m						
	total	used	free	shared	buff/cache	available
Mem:	11906	4539	3274	516	4092	8981
Swap:	0	0	0			

Private mappings. Copy-on-write

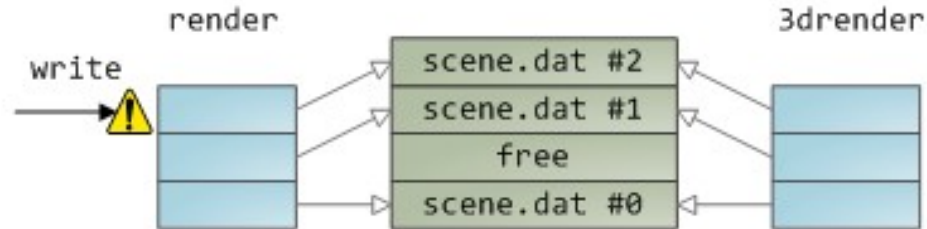
MAP_PRIVATE. The file is mapped copy-on-write, and any changes made in memory by this process are not reflected in the actual file, or in the mappings of other processes

—▷ Page table entry marked read-only
—▶ Page table entry marked read/write

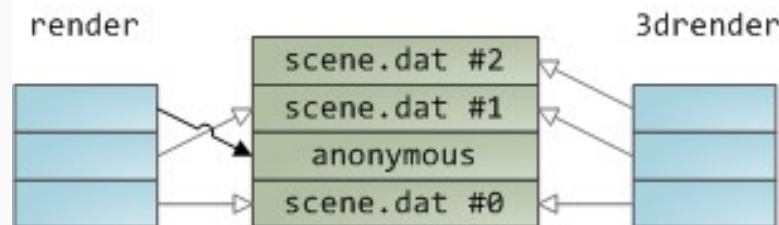
1. Two programs map scene.dat privately. Kernel deceives them and maps them both onto the page cache, but makes the PTEs read only.



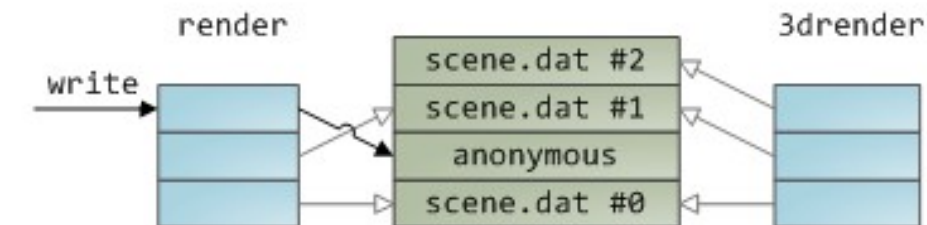
2. Render tries to write to a virtual page mapping scene.dat. Processor page faults.



3. Kernel allocates page frame, copies contents of scene.dat #2 into it, and maps the faulted page onto the new page frame.



4. Execution resumes. Neither program is aware anything happened.



Shared mappings

